


Gradient alignment at identical weights



Three fixed BP checkpoints per dataset; four disjoint training query batches each. Error bars: sample SD across checkpoints.